

CUDA Introduction

HPC - Algorithms and Applications

Alexander Pöpl
Technical University of Munich
Chair of Scientific Computing

October 31st 2018



TUM Uhrenturm

References

- D. Kirk, W. Hwu:
Programming Massively Parallel Processors, Morgan Kaufmann,
2010

Part I

Introduction to CUDA

Motivation



Currently there are two opportunities of parallelizing programs

multi-cores

- distribute the work on few strong multi-purpose processors (CPUs)
- regular supercomputers, clusters
- OpenMP, MPI

many-cores


- distribute the work on a lot of single purpose processors
- Intel MIC, BlueGene, GPUs



Difference CPU / GPU

CPU


- general purpose
- large amount of transistors for non-computational tasks
- allows out of order execution
- pipelining
- optimized for sequential tasks



CPU

GPU

- many processors dedicated to computations
- less support for branching
- well aligned data streamed through processor – data parallelism



GPU

GPU Computing – Timeline

- 80s/90s: fixed-function graphics pipelines
- 2000: programmable real-time graphics, vertex/pixel shaders
- 2003: "GPGPU", parallel programmable hardware exploited for general purposes
- 2007: GPU Computing, many-core architecture, CUDA and OpenCL programming models

Different Programming Models for GPU

CUDA

- GPU - only
- standard formed by vendor (nVidia)
- adopts new architectures fast

OpenCL



- standard formed by consortium (Khronos Group)
- platform independent (also for ATI and CPU's)
- slower development

We will use CUDA

- interface is easier to learn
- paradigms of GPU programming better understandable


GPU Architectures

NVIDIA Fermi




Fermi Streaming Multiprocessor (SM)
(source: NVIDIA – Fermi Whitepaper)

Fermi Memory Hierarchy



GPU Architectures

NVIDIA Fermi Microarchitecture



(source: NVIDIA – Fermi Whitepaper)

GPU Architectures

NVIDIA Kepler/Volta Streaming Multiprocessors



(sources: <https://devblogs.nvidia.com/parallelforall/inside-volta/> & Kepler Whitepaper)

GPU Architectures

NVIDIA Kepler Microarchitecture



(source: NVIDIA – Kepler Whitepaper)

GPU Architectures

NVIDIA Volta Microarchitecture



(source: <https://devblogs.nvidia.com/parallelforall/inside-volta/>)

Part II

CUDA Basics

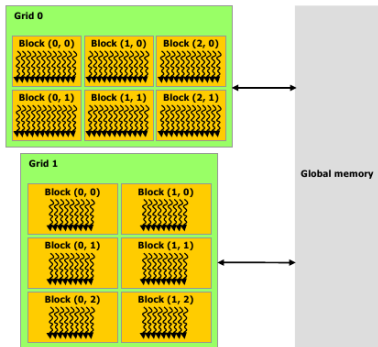

CUDA – Architecture Model

Host & Device:

- host = CPU with main memory
- device(s) = GPU/coprocessor(s) with device memory

Hardware characteristics:

- massively parallel (thousands of cores)
- lightweight threads, hardware-supported
- massive parallelism hides memory latency;
focus on data parallelism



Warps

- warp: set of 32 concurrent threads in a block
- only one (Fermi) / two (Kepler) instruction(s) per cycle per warp

Question: what happens when a warp executes a branch?

Host Memory

- slow access by PCI bus
- used as little as possible

CUDA – Programming Model

CUDA as extension of C:

- host code (program control) and device code (GPU) combined in a single C program
- device code consists of massively parallel *kernels* that are off-loaded to the GPU
- language extension for defining and calling kernels
- API function to allocate device/host memory, synchronize threads, etc.
- SIMD/SPMD (single instruction/program, multiple data)

Example: Matrix Multiplication

General Approach: PRAM program

```
for i from 1 to n do (in parallel?)  
  for k from 1 to n do (in parallel?)  
    for j from 1 to n do (in parallel?)  
      C[i,k] += A[i,j]*B[j,k]
```

Example: Matrix Multiplication

General Approach: PRAM program

```
for i from 1 to n do in parallel
  for k from 1 to n do in parallel
    for j from 1 to n do
      C[i,k] += A[i,j]*B[j,k]
```

- PRAM: executed on n^2 processors
- CUDA: n^2 CUDA threads; each thread executes one j-loop (i.e., computes one element $C[i,k]$)
- part 1: memory transfer (host→device and device→host)
- part 2: launch/execution of kernel code for j-loop

CUDA Memory transfer

Memory instructions in CUDA:

- Memory allocation:
`cudaMalloc(void** ppd, int size);`
- Memory deallocation:
`cudaFree(void* pd);`
- Copy from host to device:
`cudaMemcpy(pd, p, size, cudaMemcpyHostToDevice);`
- Copy from device to host:
`cudaMemcpy(p, pd, size, cudaMemcpyDeviceToHost);`

Matrix Multiplication – Memory Transfer

```
__host__ void matrixMult(float *A, float *B,  
float *C, int n);
```

Task:

- Write the memory transfer code for a matrix multiplication $C \leftarrow C + A \cdot B$, i.e. allocate device data, copy data between host and device and deallocate device data. All matrices have size $n \times n$.

Matrix Multiplication – Memory Transfer

```
__host__ void matrixMult(float *A, float *B,
float *C, int n) {
    int size = n*n*sizeof(float);
    float* Ad; float* Bd; float* Cd;
    cudaMalloc((void**)&Ad, size);
    cudaMalloc((void**)&Bd, size);
    cudaMalloc((void**)&Cd, size);
    cudaMemcpy(Ad,A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(Bd,B, size, cudaMemcpyHostToDevice);
    cudaMemcpy(Cd,C, size, cudaMemcpyHostToDevice);
    /* ... perform multiplication on device ... */
    cudaMemcpy(C,Cd, size, cudaMemcpyDeviceToHost);
    cudaFree(Ad); cudaFree(Bd); cudaFree(Cd);
}
```

Grids and Blocks in CUDA

Blocks:

- threads can be organized as 1D, e.g. (128,1,1), 2D, e.g. (16,16,1), or 3D, e.g. (4,8,16) blocks
- limited to 1024 (Fermi, Kepler), 512 (GT2xx) threads per block
- threads in one block are always executed in parallel
- and can use separate, shared memory

Grids:

- CUDAcc < 2.x: 2D layout. Today: 3D layout allowed
- up to $2^{32} \times 2^{32}$ (\geq Kepler) / $2^{16} \times 2^{16}$ (\leq Fermi) blocks per grid
- blocks in a grid may be executed in parallel
(but, in practice, will be scheduled to available cores)

Kernel Invocation: Grids and Blocks

```
/* ... */  
dim3 dimBlock(n,n);  
dim3 dimGrid(1,1);  
matrixMultKernel<<<dimGrid,dimBlock>>>(Ad,Bd,Cd,n);  
/* ... */
```

- threads are combined to 3D **blocks**:
→ `threadIdx.x`, `threadIdx.y`, `threadIdx.z`
- blocks are combined to 2D **grids**:
→ `blockIdx.x`, `blockIdx.y`

Matrix Multiplication – CUDA Kernel

```
__global__ void matrixMultKernel(float* Ad, float* Bd,
float* Cd, int n) {
    int i = threadIdx.x;
    int k = threadIdx.y;
    float Celem = 0;
    for(int j=0; j<n; j++) {
        float Aelem = Ad[i*n+j];
        float Belem = Bd[j*n+k];
        Celem += Aelem*Belem;
    }
    Cd[i*n+k] += Celem;
}
```


Matrix Multiplication – with Grid

```
__global__ void matrixMultKernel(float* Ad, float* Bd,  
float* Cd, int n);
```

Task:

- Extend the kernel from a single block to a grid with multiple blocks. Each block should have the size $\text{TILE_SIZE} \times \text{TILE_SIZE}$.
- Change grid and block dimensions in the kernel invocation accordingly

Matrix Multiplication – with Grid

```
__global__ void matrixMultKernel(float* Ad, float* Bd,
float* Cd, int n) {
    int i = blockIdx.x * TILE_SIZE + threadIdx.x;
    int k = blockIdx.y * TILE_SIZE + threadIdx.y;
    float Celem = 0;
    for(int j=0; j<n; j++) {
        float Aelem = Ad[i*n+j];
        float Belem = Bd[j*n+k];
        Celem += Aelem*Belem;
    }
    Cd[i*n+k] += Celem;
}
```

Matrix Multiplication – with Grid (2)

```
/* ... */  
dim3 dimBlock(TILE_SIZE,TILE_SIZE);  
dim3 dimGrid(n/TILE_SIZE,n/TILE_SIZE);  
matrixMultKernel<<<dimGrid,dimBlock>>>(Ad,Bd,Cd,n);  
/* ... */
```

- What is the optimal tile size?

Matrix Multiplication – with Grid (2)

```
/* ... */  
dim3 dimBlock(TILE_SIZE,TILE_SIZE);  
dim3 dimGrid(n/TILE_SIZE,n/TILE_SIZE);  
matrixMultKernel<<<dimGrid,dimBlock>>>(Ad,Bd,Cd,n);  
/* ... */
```

- What is the optimal tile size?
- Too small → large overhead, low performance
- Too large → block size limit, register limit
- $\text{TILE_SIZE}^2 \leq \text{max. threads per block}$
Fermi, Kepler: 1024 \Rightarrow choose $\text{TILE_SIZE} = 32$ (for now)
- If $n \bmod 32 \neq 0$: padding of matrix to match tile size
($n \leftarrow \lceil \frac{n}{32} \rceil \cdot 32$)

Matrix Multiplication – with Grid (2)

```
/* ... */  
dim3 dimBlock(TILE_SIZE,TILE_SIZE);  
dim3 dimGrid(n/TILE_SIZE,n/TILE_SIZE);  
matrixMultKernel<<<dimGrid,dimBlock>>>(Ad,Bd,Cd,n);  
/* ... */
```

- What is the optimal tile size?
- Too small → large overhead, low performance
- Too large → block size limit, register limit
- $TILE_SIZE^2 \leq \text{max. threads per block}$
Fermi, Kepler: 1024 \Rightarrow choose $TILE_SIZE = 32$ (for now)
- If $n \bmod 32 \neq 0$: padding of matrix to match tile size
($n \leftarrow \lceil \frac{n}{32} \rceil \cdot 32$)

Let's compare CPU/GPU performance!

Matrix Multiplication – with Grid (2)

```
/* ... */  
dim3 dimBlock(TILE_SIZE,TILE_SIZE);  
dim3 dimGrid(n/TILE_SIZE,n/TILE_SIZE);  
matrixMultKernel<<<dimGrid,dimBlock>>>(Ad,Bd,Cd,n);  
/* ... */
```

- What is the optimal tile size?
- Too small → large overhead, low performance
- Too large → block size limit, register limit
- $TILE_SIZE^2 \leq \text{max. threads per block}$
Fermi, Kepler: 1024 \Rightarrow choose $TILE_SIZE = 32$ (for now)
- If $n \bmod 32 \neq 0$: padding of matrix to match tile size
($n \leftarrow \lceil \frac{n}{32} \rceil \cdot 32$)

Let's compare CPU/GPU performance! \Rightarrow Okay, that sucks.

CUDA Memory

Types of **device** memory in CUDA:

- per thread: **registers** and **local memory**
(locally declared variables and arrays (local memory),
→ lifetime: kernel execution)
- per block: **shared memory**
(keyword `__shared__`, lifetime: kernel execution)
- per grid: **global memory** and **constant memory**
(keywords `__device__`, `__constant__`;
lifetime: entire application)
- vs.: CPU cache hierarchy (registers, caches, RAM)

Matrix Multiplication – Performance Estimate

Multiplication kernel:

```
for(int j=0; j<n; j++) {  
    float Aelem = Ad[i*n+j];  
    float Belem = Bd[j*n+k];  
    Celem += Aelem*Belem;  
}
```

- NVidia GT-650M stats: memory bandwidth: 64 GB/s, floating point performance: 691 GFlop/s
- two floating-point operations (multiply and add) per two floating-point variables (each 4 byte)
- thus: max. of 16 Giga `float` variables can be transferred from global memory per second
- limits performance to < 16 GFlop/s

Matrix Multiplication with Tiling

- observation: simple matrix multiplication kernel is slow (far below peak performance)
- anticipated reason: redundant access to slow global memory; performance limited by memory bandwidth between global memory and CUDA cores (each entry is loaded n times)

Remedy: **Tiling**

- switch to tile-oriented implementation (matrix multiplication on sub-blocks)
- copy matrix tiles into shared memory
- let all threads of a block work together on shared tile
- accumulate result tile back on matrix in global memory

A Note on Synchronization

Barrier-synchronization in CUDA:

```
__syncthreads();
```

- barrier for all threads within a block
- usual rules: all threads need to execute (or not) the same(!) call to `__syncthreads()`
- threads of the same block scheduled to the same hardware unit
- in contrast: no synchronization features for threads in a grid → reason: *transparent scheduling* of entire blocks

Matrix Multiplication – with Tiles

```
__global__ void matrixMultKernel(float* Ad, float* Bd,
float* Cd, int n) {
    __shared__ float Ads[TILE_SIZE][TILE_SIZE];
    __shared__ float Bds[TILE_SIZE][TILE_SIZE];
    int tx = threadIdx.x;
    int ty = threadIdx.y;
    int i = blockIdx.x * TILE_SIZE + tx;
    int k = blockIdx.y * TILE_SIZE + ty;
    /* (cont.) */
}
```

Task:

- Copy matrix tiles of A and B from global to shared memory.
- Execute matrix multiplication on shared tiles and write the result to matrix C in global memory.
- Synchronize the thread block where it is necessary.