

Tutorial: HPC - Algorithms and Applications

WS 18/19

Complete the following assignments (alone or in a group), and send *only* your source code via e-mail to poeppel@in.tum.de until Sunday, November 11th, 2018.

Worksheet 1: Matrix Multiplication in CUDA

T1.1: Basic matrix multiplication

- a) Write a simple matrix-matrix multiplication for small matrices (up to $n = 16$). The following tasks are necessary:
- Compute the amount of memory required for storing an $n \times n$ matrix with single floating point precision.
 - Allocate device memory, transfer the input matrices **A** and **B** from host to device memory and the result matrix **C** back to host memory, deallocate device memory.
 - Define grid and block size and call the device kernel. You can assume for now, that a single block is sufficient for a matrix.
 - Implement a basic matrix multiplication kernel in the function `matrixMultKernel_global`

T1.2: Increase of problem size

- a) Extend the code by allowing matrices of size $n > 16$.
- Change your grid and block size computation to handle $n \times n$ matrices for any $n > 0$. You can assume that the matrices fit into device memory.
 - Change your matrix multiplication kernel to handle the new grid and block sizes.
- b) Measure the execution time using different block sizes. Compare the execution time with the execution time of a CPU code by replacing the call to `CUDA_matrixMult` with `CPU_matrixMult`. Find the optimal block size for a matrix of size 256×256 .

H1.1: Make it run

Assignments a) and b) are mutually exclusive, complete *only one of them*.

a) Compile the exercise code on a local machine:

- Verify you have a CUDA-capable GPU and install the CUDA toolkit.
- Download and extract the exercise code from http://www5.in.tum.de/wiki/index.php/HPC_-_Algorithms_and_Applications_-_Winter_15, i.e. into `~/HPC/Exercise1`
- You might have to export some path variables:
`export PATH=$PATH:<cuda_dir>/bin`
`export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:<cuda_dir>/lib64`
- Change to the folder `cd ~/HPC/Exercise1` and type `make`, compilation should work.
- Execute the program using `./out <MATRIX-SIZE> <NUM-REPEATS>`

b) Compile the exercise code on the SCCS cluster:

- Know how to access the cluster:
 - **Host:** `pproc-be.in.tum.de`
 - **User:** `<Your TUMOnline ID>`
 - **Password:** `<Your TUMOnline Password>`
- Copy the necessary files to the cluster (using `rsync`, `scp`, or any SCP tool of your choice)
- Open an SSH connection to the cluster, then open an SSH connection into one of the GPU nodes `gpunode<1-4>`.
- Change to the folder you copied, and type `make`. Compilation should work.
- Execute the program using `./out <MATRIX-SIZE> <NUM-REPEATS>`

H1.2: Tiled matrix multiplication

a) Implement a tiled matrix multiplication kernel for improved memory performance. The tile size is defined in a preprocessor macro called `TILE_SIZE`.

- In the function `matrixMultKernel_tiled`, allocate shared memory that holds matrix tiles of size `TILE_SIZE × TILE_SIZE` for the matrices **A** and **B**.
- Fill the shared tiles with data from the matrices.
- Perform a partial matrix multiplication on the shared tiles.
- Set appropriate thread barriers in order to synchronize all threads of a block.

b) Compare the performance of CPU, basic GPU and tiled GPU matrix multiplication for a matrix of size 256×256 .

- Which implementation is the fastest?
- Are there differences for different values for `TILE_SIZE`?